

Supplementary File: Harnessing DNA Foundation Models for Cross-Species Transcription Factor Binding Site Prediction in Plant Genomes

Maryam Haghani¹, Krishna Vamsi Dhulipalla¹, Song Li^{1,2,3*}

¹Department of Computer Science, Virginia Tech, Blacksburg, VA 24061, United States of America.

²Bioinformatics, and Computational Biology, Virginia Tech, Blacksburg, VA, USA

³School of Plant and Environmental Sciences, Virginia Tech, Blacksburg, VA, USA

* Email: songli@vt.edu

S1 Dataset

S1.1 Statistics

Tables S1 - S2 summarize the dataset statistics for the one-chromosome-held-out, species-held-out, and cross-dataset experiments, respectively.

Table S1: Distribution of Sun2022 AREB/ABF transcription factor binding site samples across different chromosomes for *Arabidopsis thaliana* (*A. thaliana*).

Chromosome	# samples
Number of peak samples	14,454
Number of unique peak samples	14,374
Chromosome 1	3,787
Chromosome 2	2,232
Chromosome 3	2,807
Chromosome 4	2,297
Chromosome 5	3,251

Table S2: Dataset statistics for AREB/ABF2 binding sites in *A. thaliana* used in the cross-dataset evaluation. 735 duplicate records denote overlapping sequences between Malley2016 and Sun2022; unique samples reflect the total after deduplication.

Statistic	Malley2016	Sun2022
Number of peak samples	2,299	13,005
Unique samples after deduplication	2,299	11,995
Mean sequence length	201	201
Standard deviation	0	0
Minimum length	201	201
25% length	201	201
50% length (median)	201	201
75% length	201	201
Maximum length	201	201

Table S3: Dataset statistics of Sun2022 AREB/ABF transcription factor binding sites for *Arabidopsis thaliana* (*A. thaliana*) and *Sisymbrium irio* (*S. irio*) used in cross-species experiment.

	<i>A. thaliana</i>	<i>S. irio</i>
Number of peak samples	14,454	10,600
Number of unique peak samples	14,374	10,558
Mean	212.13	209.58
Standard Deviation (std)	5.33	3.02
Minimum (min)	208	169
25%	210	208
50% (Median)	210	209
75%	211	209
Maximum (max)	265	241

S1.2 Negative set

We generated an equal number of negative samples to match the positive set by applying dinucleotide shuffling with uShuffle [1]. This approach has been a consistent and significant practice in the field of transcription factor binding site (TFBS) prediction since 2015, particularly with the rise of deep learning models [2, 3, 4]. Dinucleotide shuffling is particularly advantageous because it creates background sequences that control for lower-order sequence biases, such as GC content, by preserving the dinucleotide frequency of the positive (binding) sequences. As a result, predictive models are compelled to capture the true sequence motifs and patterns associated with transcription factor binding, rather than relying on simple compositional differences.

S2 Implementation Details

For the motif-based approach, positively labeled sequences from the training set were provided as input to MEME. We used discriminative mode using negatively labeled (unbound) sequences to identify motifs enriched in the positive set relative to the negative set. FIMO was then used on the test set with a significance threshold parameter set to 0.05, which defines the minimum match strength required for a subsequence to be classified as a putative binding site.

All deep learning models, whether classic architectures or pretrained DNA foundation models, were trained or fine-tuned according to the following unified protocol. This unified fine-tuning/training pipeline enables a direct comparison between specialized TFBS architectures and large-scale pre-trained genomic models.

- **Model setup**

- Classic models:
 - * **DeepBind**: implemented via the MLSNET codebase [5].
 - * **BERT-TFBS**: obtained from the original GitHub repository [6].
- Pretrained DNA foundation models:
 - * Removed any existing language-modeling head.
 - * Appended a new, randomly initialized linear classification layer with two outputs (TFBS vs. non-TFBS).

- **Training Configuration**

- **Loss function**: Binary cross-entropy

- **Optimizer:** AdamW [7].
- **Regularization:** Dropout applied to mitigate overfitting [8].
- **Batching:** All sequences were padded or truncated to the dataset’s maximum input length.
- **Parameter updates:** All layers, including embeddings, were fine-tuned. The total number of trainable parameters for each model is listed in Supplementary Table S4.

Table S4: Number of trainable parameters for each model.

Model	# Parameters
DeepBind	2,194
DNABERT-2	117,070,082
BERT-TFBS	142,319,622
AgroNT	985,099,603
HyenaDNA (hyenadna-tiny-1k-seqlen)	436,354

- **Hyperparameters:** Learning rates, batch sizes, and other hyperparameters were set to each model’s default values (see Supplementary Table S5). When hyperparameter tuning was required, we employed stratified 5-fold cross-validation on the training set to maintain balanced class proportions across folds. In each iteration, one fold was used for validation, while the remaining four were used for model training.

Table S5: Default hyperparameter settings for each model. Epoch counts for fine-tuning were chosen empirically.

Model	Kernel size	Epochs	Adj. LR	Train BS	Eval BS	LR	Weight Decay
DeepBind	24	1,000*	No	64	–	1×10^{-4}	1×10^{-4}
BERT-TFBS	–	15	Yes (max 1.5×10^{-5} , min 2×10^{-6})	32	–	1.5×10^{-5}	1×10^{-2}
DNABERT-2	–	15	No	32	–	1.5×10^{-5}	1×10^{-2}
AgroNT	–	60	No	8	64	1×10^{-5}	1×10^{-1}
HyenaDNA	–	60	No	32	–	1×10^{-4}	1×10^{-1}

*Although DeepBind’s default # of epochs is 1,000, training never exceeded 100 epochs due to early stopping.

• Checkpointing

- 5% of the training data was reserved as a hold-out validation set, with stratification applied to preserve balanced class distributions across the splits.
- Validation loss was monitored each epoch, with early stopping after 10 epochs of no improvement.
- The checkpoint yielding the lowest validation loss was retained for testing.

S2.1 Evaluation Metrics

Because TFBS prediction is a binary classification task, we evaluate each model using five complementary metrics—accuracy (ACC), F1 score, Matthews correlation coefficient (MCC), ROC-AUC and PR-AUC—to provide a comprehensive assessment of classifier performance:

- **Accuracy** [9]: fraction of correctly classified samples, including TFBSs and non-TFBSs, to all the tested samples.

$$\text{Accuracy} = \frac{TP + TN}{TP + FN + TN + FP}$$

where TP, FN, TN, and FP denote the number of true positives, false negatives, true negatives, and false positives, respectively.

- **F1 score:** harmonic mean of precision (Pr) and recall (Re)

$$F1 = 2 \times \frac{\text{Pr} \times \text{Re}}{\text{Pr} + \text{Re}}$$

- **Matthews correlation coefficient (MCC)** [10]: balanced correlation coefficient, robust under imbalance

$$\text{MCC} = \frac{TP \times TN - FP \times FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

- **ROC-AUC** [11]: area under the receiver operating characteristic (ROC) curve, which plots

$$\text{TPR} = \frac{TP}{TP + FN}, \quad \text{FPR} = \frac{FP}{FP + TN}$$

over all thresholds. It is suitable for evaluating the performance of classifiers at different operating points.

- **PR-AUC** [12]: area under the precision-recall (PR) curve, trading off

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}$$

and more sensitive to class imbalance.

S3 Results

S3.1 Cross-Chromosome Evaluation on *A. thaliana* AREB/ABF1-4 Binding Sites

Figure S1 illustrates the distributions of test performance and training cost under a leave-one-chromosome-out scheme. For each of the five *A. thaliana* chromosomes, one chromosome was reserved as the test set while models were trained on the remaining four. Distributions summarize results from five models initialized with different random seeds, reporting ACC, F1 score, MCC, ROC-AUC, PR-AUC, and total training time. AgroNT consistently achieves the highest absolute values for ACC, F1, MCC, ROC-AUC, and PR-AUC. Statistical analysis using ANOVA followed by Tukey’s HSD test reveals that models sharing the same letter are not significantly different at the 0.05 significance level. This analysis shows that AgroNT, HyenaDNA, DNABERT-2, and BERT-TFBS perform comparably across all metrics and chromosome splits, often sharing the same significance group. DeepBind provides moderate improvements over the motif-based baseline, but it remains significantly weaker than transformer-based approaches, while motif-based methods are the least effective and highly variable. For computational efficiency, the contrast is stark: AgroNT is the most resource-intensive, requiring by far the longest training times, whereas HyenaDNA is exceptionally efficient, training almost 130 times faster and running inference nearly 50 times faster than AgroNT while still achieving competitive performance metrics. Overall, HyenaDNA offers the best balance between performance and computational cost.

Table S6 reports per-chromosome mean performance and training time for each model (motif-based, BERT-TFBS, DNABERT-2, DeepBind, HyenaDNA, AgroNT) under the leave-one-chromosome-out protocol. For each held-out chromosome, results are averaged over five models trained with different random seeds, highlighting variation in accuracy, correlation, AUC scores, and computational cost across distinct genomic contexts.

Table S6: Per-chromosome mean performance and timing for each model under **leave-one-chromosome-out** evaluation. Metrics include accuracy (ACC), F1 score, MCC, ROC-AUC, PR-AUC, total training time, and inference time. Each entry is the average over five models trained with different random seeds on the held-out chromosome. For each split, the top value in each column is highlighted in bold.

Inference Chromosome	Model	ACC	F1	MCC	ROC-AUC	PR-AUC	Train Time(s)	Test Time(s)
chromosome 1	Motif-based	0.780	0.764	0.570	0.804	0.796	2,800	9
	DeepBind	0.882	0.884	0.768	0.950	0.950	908	6
	BERT-TFBS	0.934	0.934	0.870	0.990	0.990	5,205	38
	DNABERT-2	0.920	0.918	0.836	0.974	0.970	5,100	37
	AgroNT	0.968	0.968	0.936	0.992	0.994	29,169	256
	HyenaDNA	0.950	0.950	0.898	0.988	0.988	215	5
chromosome 2	Motif-based	0.718	0.676	0.450	0.736	0.724	3,535	6
	DeepBind	0.866	0.864	0.728	0.934	0.938	335	4
	BERT-TFBS	0.914	0.908	0.832	0.972	0.980	5,970	22
	DNABERT-2	0.894	0.896	0.792	0.950	0.940	5,819	22
	AgroNT	0.936	0.936	0.878	0.984	0.984	33,985	151
	HyenaDNA	0.926	0.918	0.846	0.968	0.972	335	3
chromosome 3	Motif-based	0.744	0.686	0.524	0.748	0.728	3,469	7
	DeepBind	0.884	0.884	0.712	0.952	0.952	970	4
	BERT-TFBS	0.858	0.892	0.770	0.896	0.894	5,464	28
	DNABERT-2	0.912	0.914	0.830	0.968	0.960	5,568	27
	AgroNT	0.965	0.965	0.928	0.990	0.990	34,973	190
	HyenaDNA	0.940	0.942	0.880	0.980	0.982	250	3.9
chromosome 4	Motif-based	0.746	0.700	0.520	0.760	0.752	3,790	6
	DeepBind	0.882	0.884	0.768	0.950	0.950	1122	4
	BERT-TFBS	0.860	0.892	0.722	0.888	0.890	5,910	23
	DNABERT-2	0.928	0.928	0.858	0.982	0.980	5,790	23
	AgroNT	0.966	0.966	0.934	0.992	0.994	36,066	156
	HyenaDNA	0.942	0.942	0.888	0.984	0.986	254	3
chromosome 5	Motif-based	0.700	0.606	0.460	0.700	0.670	3,928	7
	DeepBind	0.886	0.886	0.774	0.952	0.954	977	5
	BERT-TFBS	0.952	0.952	0.904	0.990	0.990	5,452	33
	DNABERT-2	0.918	0.920	0.840	0.972	0.972	5,358	32
	AgroNT	0.970	0.970	0.944	0.996	0.998	34,150	220
	HyenaDNA	0.944	0.944	0.884	0.986	0.984	251	4

S3.2 Cross-Dataset Evaluation on *A. thaliana* AREB/ABF2 Binding Regions

We conducted the statistical testing between HyenaDNA and the other models using Wilcoxon test.

HyenaDNA achieved the strongest performance overall, with significantly higher accuracy, F1-score, and MCC compared to all other methods based on the Wilcoxon test (all $p = 0.010\text{--}0.011$). For ROC-AUC, HyenaDNA significantly outperformed the Motif baseline ($p = 0.009$) and DeepBind ($p = 0.007$), while differences with BERT-TFBS, DNABERT-2, and AgroNT were not statistically significant ($p = 0.408$, $p = 0.104$, and $p = 0.498$, respectively). For PR-AUC, HyenaDNA again showed significant improvements over Motif ($p = 0.009$), DeepBind ($p = 0.009$), and DNABERT-2 ($p = 0.044$), whereas differences with BERT-TFBS and AgroNT were not significant ($p = 0.071$ and $p = 0.797$). These results underscore the superiority of HyenaDNA across most performance metrics, confirming its clear statistical advantage over both motif-based and several neural sequence models in the cross-dataset setting.

In terms of training and inference time, HyenaDNA was significantly faster than all other models, including the Motif baseline ($p = 0.008$), DeepBind ($p = 0.008$), BERT-TFBS ($p = 0.008$), DNABERT-2 ($p = 0.008$), and AgroNT ($p = 0.008$).

Table S7 summarizes the mean test-set performance metrics and runtimes (averaged over five independently trained models with different random seeds) for models trained on the Malley2016

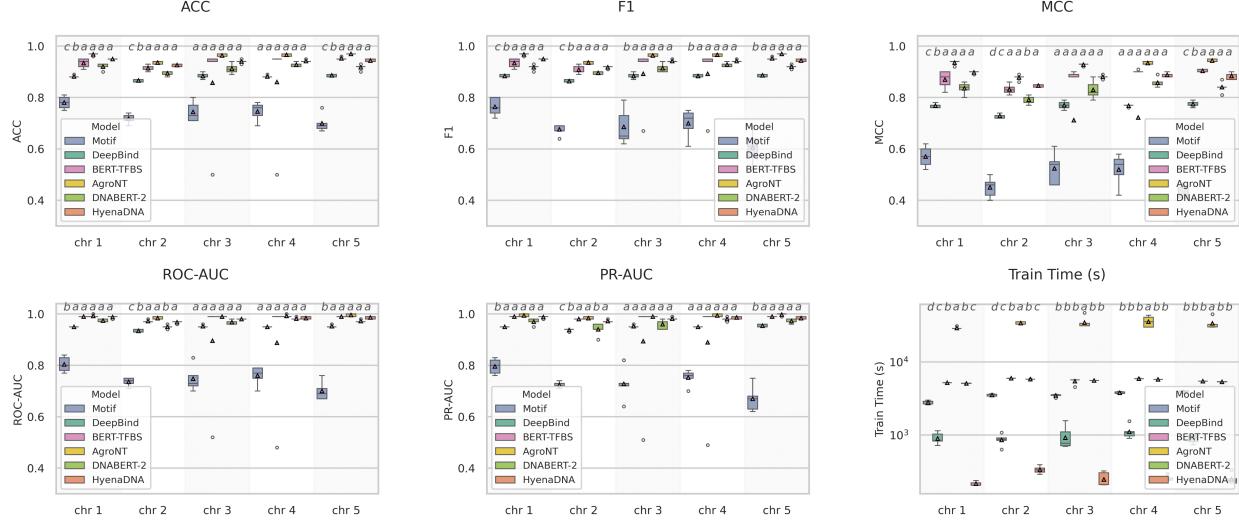


Figure S1: Per-chromosome distributions of different metrics for each model under the leave-one-chromosome-out protocol. Each marker represents the mean over five models trained with different random seeds on a single held-out chromosome. Training times are plotted on a logarithmic scale to accommodate the wide range across models. Letters on top of each box indicate statistical groupings; models sharing the same letter are not significantly different at the 0.05 significance level.

AREB/ABF2 dataset and evaluated on Sun2022 AREB/ABF2. HyenaDNA demonstrates strong generalization, substantially outperforming classical approaches such as the motif-based method, DeepBind, and BERT-TFBS, while performing on par with DNABERT-2 and AgroNT. Notably, HyenaDNA achieves this with a fraction of the computational cost—training more than $12\times$ faster than BERT-TFBS and DNABERT-2 and nearly $90\times$ faster than AgroNT, and delivering inference nearly $6\times$ faster than BERT-TFBS and DNABERT-2 and almost $40\times$ faster than AgroNT—while maintaining competitive accuracy. These results highlight HyenaDNA as an efficient and accurate choice for cross-dataset TF-binding prediction.

Table S7: Mean performance metrics and runtimes across five independently trained models with different random seeds for the **cross-dataset** evaluation, where models were trained on *Malley2016* AREB/ABF2 and evaluated on *Sun2022* AREB/ABF2. The best value in each column is shown in bold.

Model	ACC	F1	MCC	ROC-AUC	PR-AUC	Train Time (s)	Test Time (s)
Motif-based	0.786	0.754	0.590	0.790	0.774	3,099	24
DeepBind	0.804	0.766	0.638	0.908	0.920	162	16
BERT-TFBS	0.788	0.738	0.630	0.910	0.924	835	99
DNABERT-2	0.860	0.844	0.734	0.908	0.932	846	96
HyenaDNA	0.884	0.876	0.784	0.942	0.958	76	17
AgroNT	0.840	0.810	0.716	0.944	0.958	6,808	662

S3.3 Cross-Species Transferability of AREB/ABF1-4 Binding Sites

Table S7 reports the mean test-set performance metrics and runtimes for cross-species evaluation between *A. thaliana* and *S. irio*, averaged over five independently trained models with different random seeds. HyenaDNA closely matches the predictive accuracy of leading transformer-based (e.g. DNABERT-2) and convolutional (e.g. AgroNT) baselines while training an order of magnitude faster than DNABERT-2 and over $100\times$ faster than AgroNT. Inference time is likewise reduced by a similar factor. This combination of state-of-the-art generalization and lightweight computation

makes HyenaDNA particularly well suited for rapid, large-scale transfer learning across diverse plant genomes.

Table S8: Mean performance metrics of each model for predicting AREB/ABF transcription factor binding regions in a **cross-species** setting using *A. thaliana* and *S. irio*. Metrics are averaged across five independently trained models with different random seeds. Both directions—training on one species and testing on the other—are reported. For each direction, the top value in each column is highlighted in bold.

Direction	Model	ACC	F1	MCC	ROC-AUC	PR-AUC	Train Time(s)	Test Time(s)
train: <i>A. thaliana</i> test: <i>S. irio</i> .	Motif-based	0.850	0.834	0.720	0.858	0.852	3,037	22
	DeepBind	0.910	0.910	0.824	0.964	0.966	974	15
	BERT-TFBS	0.956	0.956	0.910	0.990	0.990	7,066	105
	DNABERT-2	0.934	0.936	0.870	0.974	0.974	6,905	104
	AgroNT	0.966	0.966	0.932	0.990	0.990	38,344	717
train: <i>S. irio</i> . test: <i>A. thaliana</i>	HyenaDNA	0.948	0.948	0.900	0.986	0.988	302	14
	Motif-based	0.708	0.708	0.426	0.768	0.754	3,450	31
	DeepBind	0.876	0.870	0.754	0.944	0.946	724	21
	BERT-TFBS	0.926	0.922	0.852	0.968	0.978	5,179	143
	DNABERT-2	0.918	0.912	0.830	0.964	0.964	5,508	140
	AgroNT	0.946	0.946	0.898	0.986	0.990	30,531	977
	HyenaDNA	0.928	0.926	0.854	0.972	0.974	213	20

S4 Ablation of HyenaDNA Freezing Regimes

To assess how different parameter-freezing strategies affect predictive performance and computational cost of HyenaDNA, we trained the model under three regimes: freezing the entire backbone (**backbone**), freezing only the backbone layers (**backbone.layers**), and no freezing (**none**). Table S9 lists the trainable parameters under these freezing regimes:

Table S9: Trainable parameter counts for HyenaDNA under different freezing strategies.

Frozen Components	Trainable Parameters
Backbone	258
Backbone.layers	2,562
None	436,354

S4.1 Cross-Chromosome Evaluation

To assess the impact of parameter freezing on HyenaDNA’s capacity to capture chromosome-specific features within a single species, we trained five independent models (with different random seeds) on a cross-chromosome dataset (Sun2022 AREB/ABF transcription factor binding site samples in *A. thaliana*), where each chromosome was held out in turn for testing. Figure S2 presents the distributions average Accuracy (ACC), F1 score, Matthews correlation coefficient (MCC), ROC-AUC, and PR-AUC across five leave-one-chromosome-out experiments for three finetuning strategies. Models trained without freezing (pink) achieve the best performance across all metrics, demonstrating strong within-species generalization. Freezing only the backbone layers (green) results in a modest reduction, suggesting that key representations are retained while still allowing adaptation. In contrast, freezing the entire backbone (purple) substantially impairs performance, highlighting the importance of updating at least part of the pretrained parameters even in within-species settings.

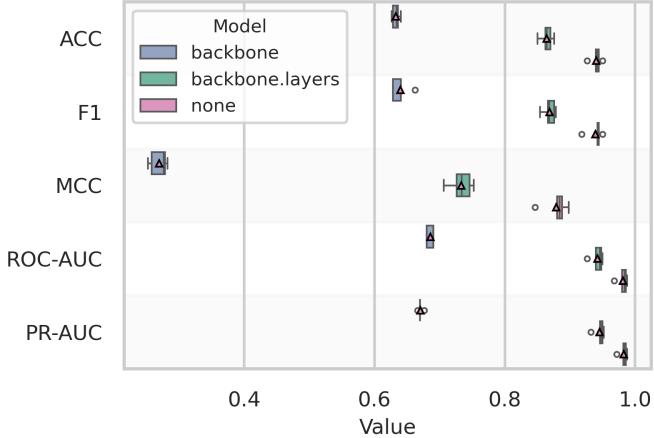


Figure S2: Distribution of test-set performance metrics for three parameter-freezing strategies in the HyenaDNA model, assessed using **five cross-chromosome hold-out** runs. For each cross-chromosome hold-out, five models trained with different random seeds and test set were evaluated. Boxplots show the distribution of mean Accuracy (ACC), F1 score, Matthews correlation coefficient (MCC), area under the ROC curve (ROC-AUC), and area under the precision-recall curve (PR-AUC) across the cross-chromosome hold-outs. Purple, green, and pink represent freezing the entire backbone, freezing only the backbone layers, and no freezing, respectively. Mean results for individual runs across the five runs are shown as circles, while averages are indicated with triangles.

S4.2 Cross-Dataset Evaluation

We trained the HyenaDNA model under the three regimes using cross-dataset experiment in which models trained on the **Malley2016** AREB/ABF2 dataset were tested on the independent **Sun2022** *A. thaliana* AREB/ABF2 dataset. We compared three finetuning regimes: (i) freezing the entire pretrained backbone, (ii) freezing only the backbone layers, and (iii) full finetuning (no freezing). Performance was measured over five independent models (with different random seeds), reporting Accuracy (ACC), F1 score, Matthews correlation coefficient (MCC), ROC-AUC and PR-AUC for each run (Figure S3).

When all parameters were trainable (pink), HyenaDNA adapted most effectively to the new dataset, achieving the highest mean scores across all metrics. Partial freezing of backbone layers (green) yielded intermediate performance, indicating some preservation of pretrained features but reduced flexibility to dataset-specific signal. Fully freezing the backbone (purple) severely degraded cross-dataset generalization, specially in terms of MCC, underscoring that without any parameter updates HyenaDNA cannot accommodate differences between protocols. These results demonstrate that model finetuning is essential for high-fidelity prediction when transferring between distinct experimental datasets.

S4.3 Cross-Species Evaluation

We performed two leave-one-species-out experiments on *A. thaliana* and *S. irio* species using **Sun2022** AREB/ABF1-4 peak regions under the three regimes (Figure S4). In each case, five different seeds were used for training the model on one species and testing on the other. When all parameters were unfrozen (pink), the model attained excellent transfer performance regardless of training direction. By contrast, freezing the entire backbone (purple) almost completely abrogated cross-species prediction, showing that without any parameter updates the model cannot adapt to species-specific

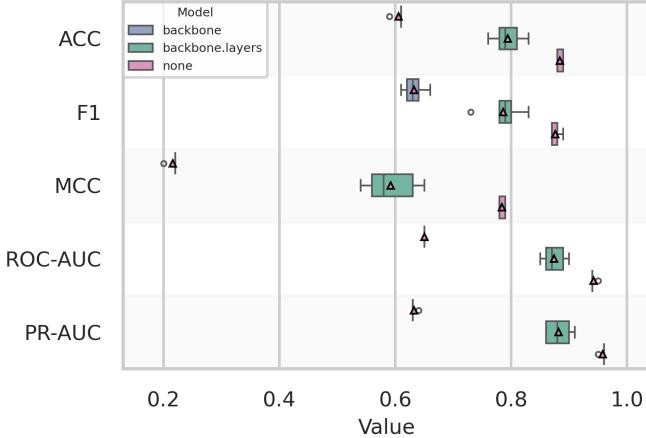


Figure S3: Distribution of test-set performance (ACC, F1, MCC, ROC-AUC, PR-AUC) per trained model for finetuning strategies of HyenaDNA in the **cross-dataset** evaluation (models trained on **Malley2016 AREB/ABF2**, tested on **Sun2022 *A. thaliana* AREB/ABF2**). Purple, green and pink correspond to freezing the entire backbone, freezing only backbone layers, and no freezing, respectively. Individual results are overlaid as circles and means as triangles

signals. Selectively freezing only the backbone layers (green) provided an intermediate solution: it recovered most of the pretrained knowledge while still allowing adaptation. These results underscore that end-to-end finetuning is essential for robust cross-species generalization, and that partial backbone freezing offers a favourable trade-off between leveraging pretrained representations and adapting to new genomic contexts.

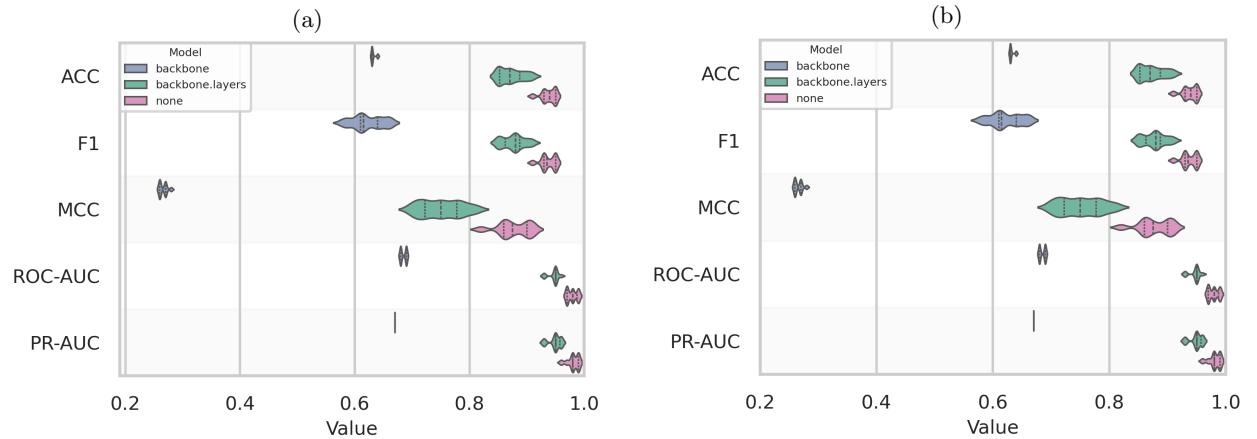


Figure S4: Distribution of test-set performance metrics for three parameter-freezing strategies in the HyenaDNA model under **cross-species** evaluation. Each distribution summarizes results from five independently trained models (with different random seeds), where models are evaluated on the held-out test species. (a): trained on *A. thaliana*, tested on *S. irio*; (b): trained on *S. irio*, tested on *A. thaliana*. Purple, green and pink correspond to freezing the entire backbone, freezing only backbone layers, and no freezing, respectively. Individual results are overlaid as circles and means as triangles.

References

- [1] M. Jiang, J. Anderson, J. Gillespie, and M. Mayne. uShuffle: a useful tool for shuffling biological sequences while preserving the k-let counts. *BMC bioinformatics*, 9(1):192, 2008.
- [2] Babak Alipanahi, Andrew Delong, Matthew T Weirauch, and Brendan J Frey. Predicting the sequence specificities of DNA-and RNA-binding proteins by deep learning. *Nature Biotechnology*, 33(8):831–838, 2015.
- [3] A. Patel, A. Singhal, A. Wang, A. Pampari, M. Kasowski, and A. Kundaje. DART-Eval: A comprehensive DNA language model evaluation benchmark on regulatory DNA. *Advances in Neural Information Processing Systems*, 37:62024–62061, 2024.
- [4] M. Tognon, A. Kumbara, A. Betti, L. Ruggeri, and R. Giugno. Benchmarking transcription factor binding site prediction models: a comparative analysis on synthetic and biological data. *Briefings in Bioinformatics*, 26(4):bbaf363, 2025.
- [5] Yuchuan Zhang, Zhikang Wang, Fang Ge, Xiaoyu Wang, Yiwen Zhang, Shanshan Li, Yuming Guo, Jiangning Song, and Dong-Jun Yu. MLSNet: a deep learning model for predicting transcription factor binding sites. *Briefings in Bioinformatics*, 25(6):bbae489, 2024.
- [6] Kai Wang, Xuan Zeng, Jingwen Zhou, Fei Liu, Xiaoli Luan, and Xinglong Wang. BERT-TFBS: a novel BERT-based model for predicting transcription factor binding sites by transfer learning. *Briefings in Bioinformatics*, 25(3):bbae195, 2024.
- [7] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*, 2017.
- [8] Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research*, 15(1):1929–1958, 2014.
- [9] John A Swets. Measuring the accuracy of diagnostic systems. *Science*, 240(4857):1285–1293, 1988.
- [10] Davide Chicco and Giuseppe Jurman. The advantages of the matthews correlation coefficient (MCC) over F1 score and accuracy in binary classification evaluation. *BMC genomics*, 21:1–13, 2020.
- [11] John Muschelli III. ROC and AUC with a binary predictor: a potentially misleading metric. *Journal of classification*, 37(3):696–708, 2020.
- [12] Helen R Sofaer, Jennifer A Hoeting, and Catherine S Jarnevich. The area under the precision-recall curve as a performance metric for rare binary events. *Methods in Ecology and Evolution*, 10(4):565–577, 2019.